

## L2. NoSQL<sup>1</sup>

### S1. Principles

Dr A Boronat

---


<sup>1</sup>Materials based on [1, ch 1-3]

# Table of Contents

- ① Motivation
- ② Outline
- ③ Schemaless
- ④ Integration
- ⑤ Big Data

## Up to now...relational databases

- Standard model (relational) and query language (SQL): different vendors' SQL dialects are similar
  - **record**: set of name-value pairs, represented as a **row**
  - **relation**: set of records, represented as a **table** (with columns and rows)
- **ACID transactions**: Atomicity, Consistency, Isolation, Durability
  - **Atomicity**: all or nothing
  - **Consistency**: any transaction will bring the database from one valid state to another
  - **Isolation**: the concurrent execution of transactions results in a system state that would be obtained if transactions were executed serially
  - **Durability**: the results need to be stored permanently even in the event of app crashes
- Used as integration mechanism between applications
  - multiple apps store their data in a single db
  - **views** help defining domain-oriented queries



## Do they cover all current needs?

- Simplicity of data model
  - a relational tuple cannot contain any structure
  - in-memory data structures can be much richer 😞
- Impedance mismatch
  - mapping data between in-memory data structures and a relational db
  - ORM frameworks (Hibernate, JPA-like) help in this task
  - however, once we forget the db structure, our query performance may worsen dramatically 😞

## Do they cover all current needs?

- Integration databases are complex
  - each app is developed by a different team and it may only use part of the database
  - someone needs to ensure the consistency of the whole database 😞
- Big data and traffic
  - scaling up (supercomputer) is expensive
  - scaling out (cluster) may be reasonably affordable
    - cheaper (and also less reliable) machines working together
    - relational databases not designed for clusters 😞

# NoSQL: a new neighbourhood

- Schemaless: more flexibility
  - no need to map structured data
  - rapid development
- Shift from integration dbs to application dbs
  - SOA and micro-services as integration paradigm
  - interoperability concerns at the interface of the application
- Cluster-oriented
  - load balancing
  - ability to trade off traditional consistency for other useful properties (e.g. availability)
- Do not use the relational model
- Tend to be open source



# Reasons

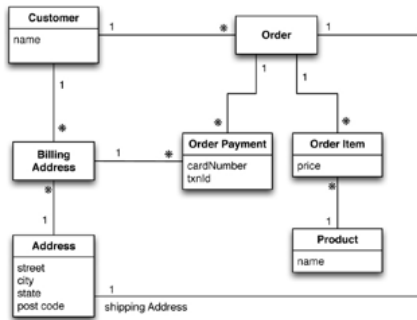
- App dev productivity
  - mapping data between in-memory data structures and a relational db
  - a NoSQL DB simplifies this interaction (less code)
- Large-scale data
  - relational dbs are designed to run on a single machine
  - NoSQL DBs are designed explicitly to run on clusters:
    - better fit for big data scenarios
    - easier to manage computing loads

## Technology Outline: aggregate-based stores

- **Aggregate**: collection of data that we interact with as a unit
  - manipulation of multiple aggregates should be dealt with in the app logic
- Form the boundaries for ACID operations with the database
- Aggregation is a physical property, it helps running on a cluster



# (Relational) data model and instance



Customer	
Id	Name
1	Martin

Order		
Id	CustomerId	ShippingAddressId
99	1	77

Product	
Id	Name
27	NoSQL Distilled

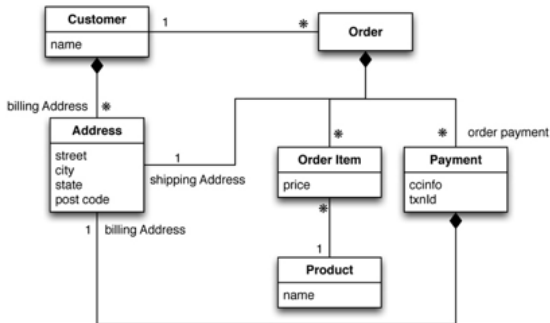
BillingAddress		
Id	CustomerId	AddressId
55	1	77

OrderItem			
Id	OrderId	ProductId	Price
100	99	27	32.45

Address	
Id	City
77	Chicago

OrderPayment				
Id	OrderId	CardNumber	BillingAddressId	txnId
33	99	1000-1000	55	abel1f879rft

# Aggregate data model and instance



```
// in customers
{
  "id":1,
  "name":"Martin",
  "billingAddress":[{"city":"Chicago"}]
}

// in orders
{
  "id":99,
  "customerId":1,
  "orderItems":[
    {
      "productId":27,
      "price": 32.45,
      "productName": "NoSQL Distilled"
    }
  ],
  "shippingAddress":[{"city":"Chicago"}]
  "orderPayment":[
    {
      "ccinfo":"1000-1000-1000-1000",
      "txnId":"abelif879rft",
      "billingAddress": {"city": "Chicago"}
    }
  ],
}
```

## Technology Outline: aggregate-based stores

- **Key-value stores:**
  - a db consists of several aggregates, each of which has a key or ID used to look up the data
  - the aggregate is opaque to the db  
Examples: Redis, Riak
- **Document stores:**
  - query based on the internal structure of the document, which can be a key
  - db is able to see the structure of the aggregate
  - can retrieve part of the aggregate rather than the whole thing  
Examples: CouchDB, MongoDB
- **Column-family:**
  - divide the aggregate into column families
  - aggregate normally computed using map-reduce algorithm  
Examples: HBase, Apache Cassandra

## Technology Outline: graph databases

- Reject SQL model
  - small records with complex interconnections
  - traversal of links is very efficient
- Suitable when
  - more queries than inserts (analytics)
  - data with complex relationships: social networks, product preferences, eligibility rules
- More likely to run on a single server
- ACID transactions need to cover multiple nodes and edges to cover consistency
- Example: Neo4J
  - allows us to attach Java objects as properties (features) to nodes and edges in a schemaless fashion


## Working without Schemas

- NoSQL is untyped: advantages
  - **flexibility for prototyping** within an aggregate
  - **nonuniform data**: each record has a different set of fields (one row per table)
  - reporting of structured data
- Working with schemas introduces discipline in programming
  - lacking a schema may introduce errors that can only be found at run time
  - the code needs to be inspected to find bugs
- The schema is **implicit**
  - the schema is shifted to the code
  - db remains ignorant of schema
    - worse performance
    - data integrity cannot be guaranteed
  - migration as cumbersome as in relational dbs in some cases
- Queries that select rows cannot be defined (use indexes instead)

## Integration with NoSQL: approaches


### through a dedicated application

- encapsulate all db interaction within a single app
- integrate it with other apps using web services



### aggregates linked to specific apps

- db is partitioned according to apps





# BIG DATA



**VOLUME**

DATA SIZE



**VELOCITY**

SPEED OF CHANGE



**VARIETY**

DIFFERENT FORMS  
OF DATA SOURCES



**VERACITY**

UNCERTAINTY OF  
DATA

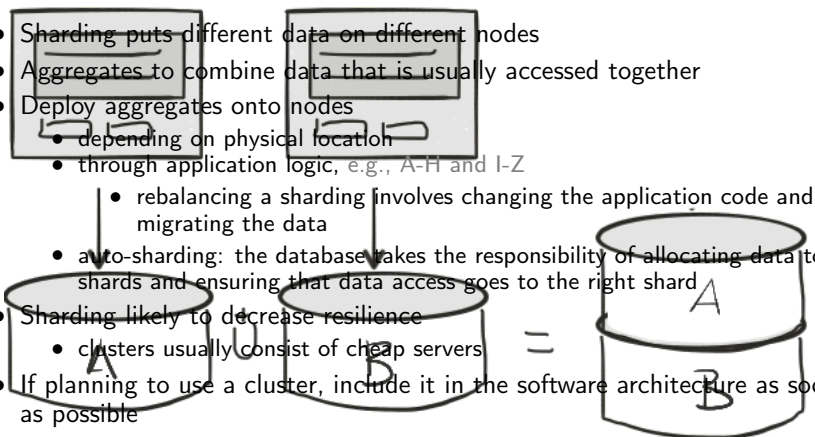
## Scaling techniques

DATA  
is the  
new  
OIL

- To deal with
  - big data: data from activity tracking
  - high number of users / social networks
- Physical approaches
  - scaling up: add resources to a server (supercomputer)
    - Ex: [Sunway TaihuLight](#): fastest supercomputer as of June 2017, ~ £207m
  - scaling out: run the db on a cluster of servers
    - less reliable
    - less powerful
    - affordable
    - IaaS databases: MongoLab
    - ALICE HPC cluster, ~£2m
- Logical approaches
  - sharding: partition database
  - replication: make several replicas available




# Sharding

- Sharding puts different data on different nodes
  - Aggregates to combine data that is usually accessed together
    - depending on physical location
    - through application logic, e.g., A-H and I-Z
  - Deploy aggregates onto nodes
    - rebalancing a sharding involves changing the application code and migrating the data
    - auto-sharding: the database takes the responsibility of allocating data to shards and ensuring that data access goes to the right shard
  - Sharding likely to decrease resilience
    - clusters usually consist of cheap servers
  - If planning to use a cluster, include it in the software architecture as soon as possible
- 

# Replication


- Replicas are deployed onto different nodes
- **Master-slave replication**
  - master: responsible for processing any updates to data
    - reduces the chance of update (write) conflicts
  - slave: read-only
    - can be used as hot-backup
    - scale out by adding slave nodes
- **advantages**
  - balances read traffic
  - read resilience: should the master fail, slave can still handle read requests
- **problems**
  - inconsistencies when updates have not been propagated to slaves
  - resilience: the master is a bottleneck and a single point of failure for writes



# Replication

- **Peer-to-Peer Replication**

- all the replicas have equal weight and they can all accept writes
- **advantages**
  - more resilient (when a node fails the others still work)
- **problems**
  - consistency: write-write conflicts (update the same record at the same time)
  - replicas coordinate to ensure we avoid a conflict (more traffic)
  - cope with an inconsistent write by allowing them and merging inconsistent writes



# What to remember

- Welcome to polyglot persistence:
  - Aggregate-oriented databases facilitate the use of clusters
  - Graph databases suitable for complex relationship structures
  - Schemaless databases enable rapid prototyping but there is an implicit schema
- Scaling techniques
  - Sharding distributes different data across multiple servers
  - Replication copies data across multiple servers

## References



Pramod J. Sadalage and Martin Fowler.

*NoSQL Distilled: A Brief Guide to the Emerging World of Polyglot Persistence.*

Addison-Wesley Professional, 1st edition, 2012.